

# lab09PCA

Ziyuan\_Han

10/27/2021

#Get the dataset into the project

```
fna.data <- "WisconsinCancer.csv"
wisc.df <- read.csv(fna.data, row.names=1)
head(wisc.df)
```

```
##      diagnosis radius_mean texture_mean perimeter_mean area_mean
## 842302         M      17.99       10.38         122.80      1001.0
## 842517         M      20.57       17.77         132.90      1326.0
## 84300903        M      19.69       21.25         130.00      1203.0
## 84348301         M      11.42       20.38          77.58       386.1
## 84358402         M      20.29       14.34         135.10      1297.0
## 843786         M      12.45       15.70          82.57       477.1
##      smoothness_mean compactness_mean concavity_mean concave.points_mean
## 842302          0.11840          0.27760          0.3001          0.14710
## 842517          0.08474          0.07864          0.0869          0.07017
## 84300903         0.10960          0.15990          0.1974          0.12790
## 84348301         0.14250          0.28390          0.2414          0.10520
## 84358402         0.10030          0.13280          0.1980          0.10430
## 843786          0.12780          0.17000          0.1578          0.08089
##      symmetry_mean fractal_dimension_mean radius_se texture_se perimeter_se
## 842302          0.2419          0.07871      1.0950      0.9053          8.589
## 842517          0.1812          0.05667      0.5435      0.7339          3.398
## 84300903         0.2069          0.05999      0.7456      0.7869          4.585
## 84348301         0.2597          0.09744      0.4956      1.1560          3.445
## 84358402         0.1809          0.05883      0.7572      0.7813          5.438
## 843786          0.2087          0.07613      0.3345      0.8902          2.217
##      area_se smoothness_se compactness_se concavity_se concave.points_se
## 842302      153.40      0.006399      0.04904      0.05373      0.01587
## 842517       74.08      0.005225      0.01308      0.01860      0.01340
## 84300903      94.03      0.006150      0.04006      0.03832      0.02058
## 84348301      27.23      0.009110      0.07458      0.05661      0.01867
## 84358402      94.44      0.011490      0.02461      0.05688      0.01885
## 843786      27.19      0.007510      0.03345      0.03672      0.01137
##      symmetry_se fractal_dimension_se radius_worst texture_worst
## 842302      0.03003          0.006193          25.38          17.33
## 842517      0.01389          0.003532          24.99          23.41
## 84300903      0.02250          0.004571          23.57          25.53
## 84348301      0.05963          0.009208          14.91          26.50
## 84358402      0.01756          0.005115          22.54          16.67
## 843786      0.02165          0.005082          15.47          23.75
```

```
##      perimeter_worst area_worst smoothness_worst compactness_worst
## 842302      184.60      2019.0          0.1622          0.6656
## 842517      158.80      1956.0          0.1238          0.1866
## 84300903     152.50      1709.0          0.1444          0.4245
## 84348301       98.87       567.7          0.2098          0.8663
## 84358402     152.20      1575.0          0.1374          0.2050
## 843786      103.40       741.6          0.1791          0.5249
##      concavity_worst concave.points_worst symmetry_worst
## 842302      0.7119          0.2654          0.4601
## 842517      0.2416          0.1860          0.2750
## 84300903     0.4504          0.2430          0.3613
## 84348301     0.6869          0.2575          0.6638
## 84358402     0.4000          0.1625          0.2364
## 843786      0.5355          0.1741          0.3985
##      fractal_dimension_worst
## 842302      0.11890
## 842517      0.08902
## 84300903     0.08758
## 84348301     0.17300
## 84358402     0.07678
## 843786      0.12440
```

#visualize the data

```
wisc.data <- wisc.df[, -1]
# Create diagnosis vector for later
diagnosis <- as.factor(wisc.df$diagnosis)
```

#Q1. How many observations are in this dataset? 569 observations

```
nrow(wisc.df)
```

```
## [1] 569
```

Q2. How many of the observations have a malignant diagnosis? 212 observations have the malignant diagnosis

```
length(which(diagnosis=="M"))
```

```
## [1] 212
```

Q3. How many variables/features in the data are suffixed with \_mean? 10

```
colname<-colnames(wisc.df)
length(grep("_mean", colname))
```

```
## [1] 10
```

#PCA analysis ##Check column means and standard deviations

```
colMeans(wisc.data)
```

```
##          radius_mean      texture_mean      perimeter_mean
##      1.412729e+01      1.928965e+01      9.196903e+01
##          area_mean      smoothness_mean      compactness_mean
##      6.548891e+02      9.636028e-02      1.043410e-01
##      concavity_mean      concave.points_mean      symmetry_mean
##      8.879932e-02      4.891915e-02      1.811619e-01
##      fractal_dimension_mean      radius_se      texture_se
##      6.279761e-02      4.051721e-01      1.216853e+00
##      perimeter_se      area_se      smoothness_se
##      2.866059e+00      4.033708e+01      7.040979e-03
##      compactness_se      concavity_se      concave.points_se
##      2.547814e-02      3.189372e-02      1.179614e-02
##      symmetry_se      fractal_dimension_se      radius_worst
##      2.054230e-02      3.794904e-03      1.626919e+01
##      texture_worst      perimeter_worst      area_worst
##      2.567722e+01      1.072612e+02      8.805831e+02
##      smoothness_worst      compactness_worst      concavity_worst
##      1.323686e-01      2.542650e-01      2.721885e-01
##      concave.points_worst      symmetry_worst      fractal_dimension_worst
##      1.146062e-01      2.900756e-01      8.394582e-02
```

```
apply(wisc.data,2,sd)
```


```
##          radius_mean      texture_mean      perimeter_mean
##      3.524049e+00      4.301036e+00      2.429898e+01
##          area_mean      smoothness_mean      compactness_mean
##      3.519141e+02      1.406413e-02      5.281276e-02
##      concavity_mean      concave.points_mean      symmetry_mean
##      7.971981e-02      3.880284e-02      2.741428e-02
##      fractal_dimension_mean      radius_se      texture_se
##      7.060363e-03      2.773127e-01      5.516484e-01
##      perimeter_se      area_se      smoothness_se
##      2.021855e+00      4.549101e+01      3.002518e-03
##      compactness_se      concavity_se      concave.points_se
##      1.790818e-02      3.018606e-02      6.170285e-03
##      symmetry_se      fractal_dimension_se      radius_worst
##      8.266372e-03      2.646071e-03      4.833242e+00
##      texture_worst      perimeter_worst      area_worst
##      6.146258e+00      3.360254e+01      5.693570e+02
##      smoothness_worst      compactness_worst      concavity_worst
##      2.283243e-02      1.573365e-01      2.086243e-01
##      concave.points_worst      symmetry_worst      fractal_dimension_worst
##      6.573234e-02      6.186747e-02      1.806127e-02
```

Perform PCA on wisc.data by completing the following code

```
wisc.pr <- prcomp(wisc.data,scale.=TRUE)
summary(wisc.pr)
```

```
## Importance of components:
##          PC1    PC2    PC3    PC4    PC5    PC6    PC7
## Standard deviation  3.6444 2.3857 1.67867 1.40735 1.28403 1.09880 0.82172
## Proportion of Variance 0.4427 0.1897 0.09393 0.06602 0.05496 0.04025 0.02251
## Cumulative Proportion 0.4427 0.6324 0.72636 0.79239 0.84734 0.88759 0.91010
##          PC8    PC9    PC10   PC11   PC12   PC13   PC14
## Standard deviation  0.69037 0.6457 0.59219 0.5421 0.51104 0.49128 0.39624
## Proportion of Variance 0.01589 0.0139 0.01169 0.0098 0.00871 0.00805 0.00523
## Cumulative Proportion 0.92598 0.9399 0.95157 0.9614 0.97007 0.97812 0.98335
##          PC15   PC16   PC17   PC18   PC19   PC20   PC21
## Standard deviation  0.30681 0.28260 0.24372 0.22939 0.22244 0.17652 0.1731
## Proportion of Variance 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104 0.0010
## Cumulative Proportion 0.98649 0.98915 0.99113 0.99288 0.99453 0.99557 0.9966
##          PC22   PC23   PC24   PC25   PC26   PC27   PC28
## Standard deviation  0.16565 0.15602 0.1344 0.12442 0.09043 0.08307 0.03987
## Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005
## Cumulative Proportion 0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
##          PC29   PC30
## Standard deviation  0.02736 0.01153
## Proportion of Variance 0.00002 0.00000
## Cumulative Proportion 1.00000 1.00000
```

```
plot(wisc.pr)
```



#Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)? 44.27% of the variance is captured


#Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data? 3 PCs are required to describe 72.6% of the original variance

#Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data? 7 PCs are required to describe at least 90% of the original variance in the data

## Interpreting PCA results


#Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why? It has too many labels and looks like messy. It's difficult to understand because the annotation is not clear.

```
library(ggplot2)
biplot(wisc.pr)
```




# Scatter plot observations by components 1 and 2

```
plot(wisc.pr$x, col = diagnosis,
     xlab = "PC1", ylab = "PC2")
```



#Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots? the PC1 and PC2 plot look similar to PC1 and PC3 plot, which indicate the first two PCs can explain most variance in the data

```
plot(wisc.pr$x[,1:3], col = diagnosis,  
     xlab = "PC1", ylab = "PC3")
```

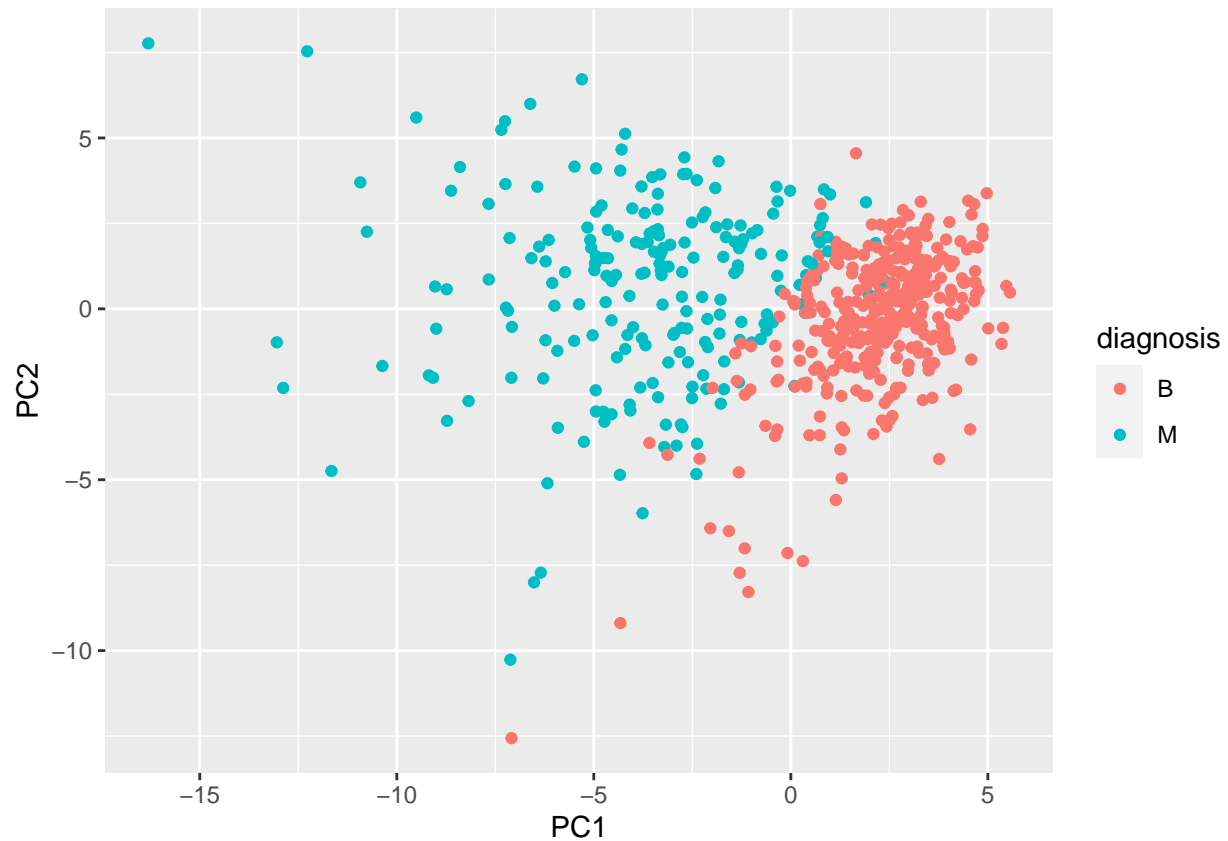


```
# Create a data.frame for ggplot
```

```
df <- as.data.frame(wisc.pr$x)
df$diagnosis <- diagnosis
```

```
# Load the ggplot2 package
library(ggplot2)
```

```
# Make a scatter plot colored by diagnosis
ggplot(df) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()
```



#Variance Explained # Calculate variance of each component

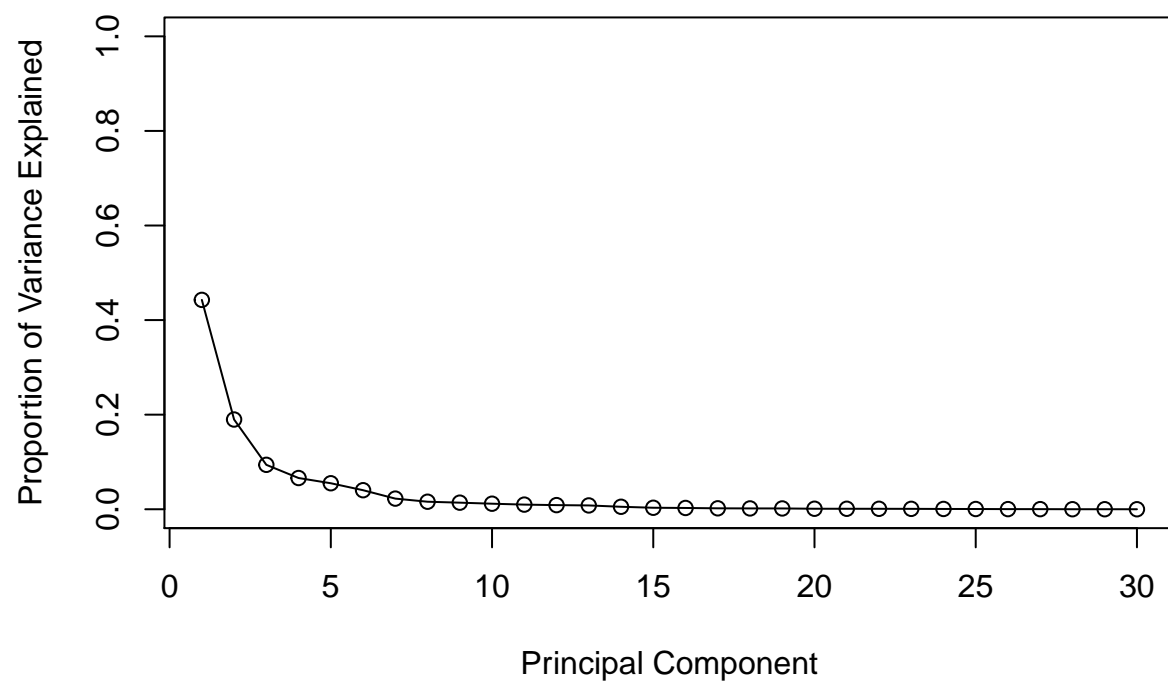
```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

```
## [1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```

```
# Variance explained by each principal component: pve
pve <- pr.var / sum(pr.var)

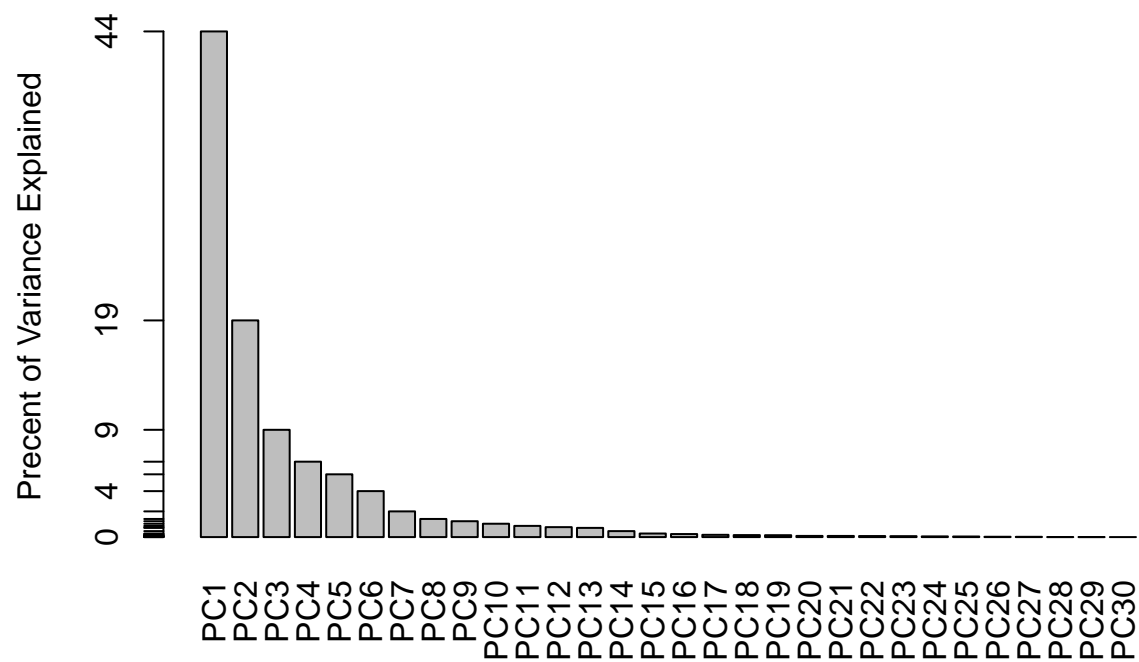
# Plot variance explained for each principal component
plot(pve, xlab = "Principal Component",
     ylab = "Proportion of Variance Explained",
     ylim = c(0, 1), type = "o")
```





# Alternative scree plot of the same data, note data driven y-axis

```
barplot(pve, ylab = "Precent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```

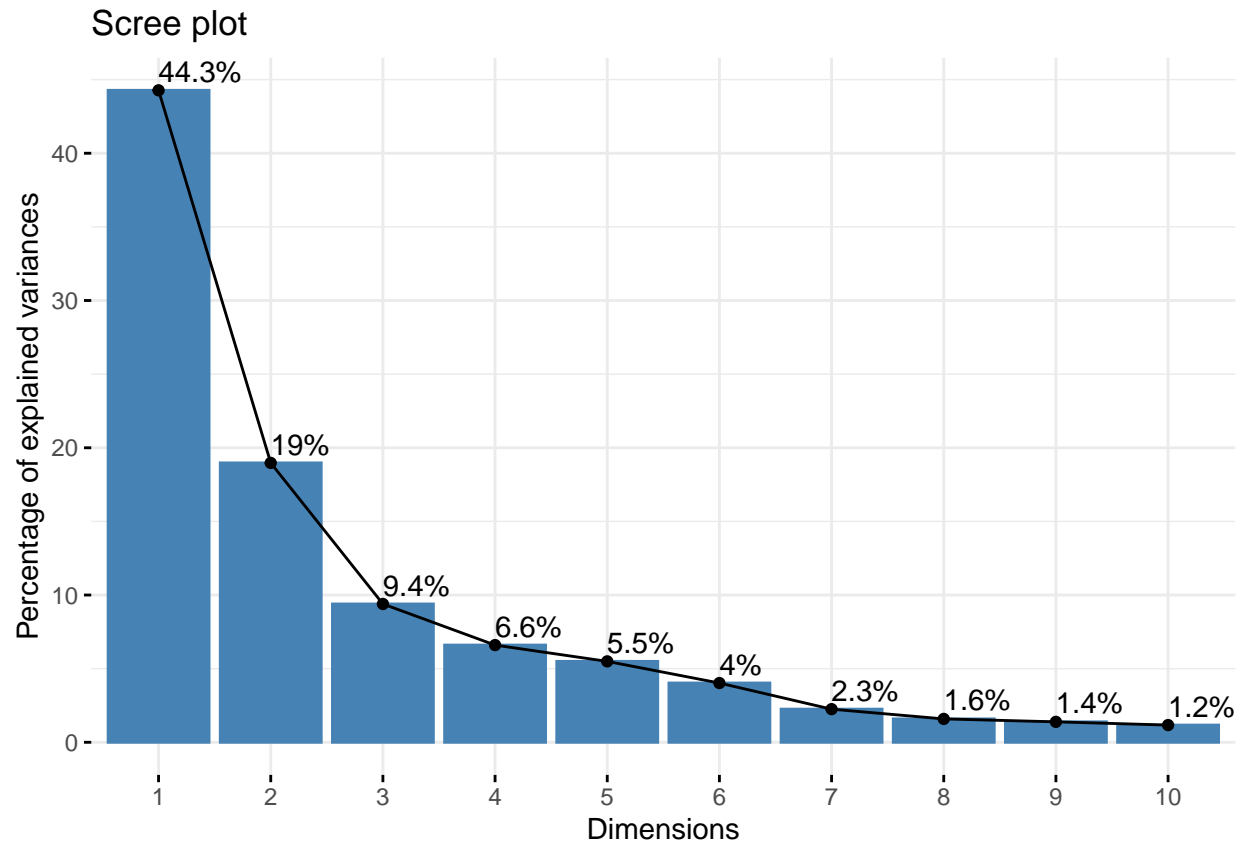


```
## ggplot based graph
```

```
#install.packages("factoextra")
library(factoextra)
```

```
## Welcome! Want to learn more? See two factoextra-related books at https://goo.gl/ve3WBa
```

```
fviz_eig(wisc.pr, addlabels = TRUE)
```



#Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`? -0.26085376

```
wisc.pr$rotation[,1]
```

```
##      radius_mean      texture_mean      perimeter_mean
##      -0.21890244      -0.10372458      -0.22753729
##      area_mean      smoothness_mean      compactness_mean
##      -0.22099499      -0.14258969      -0.23928535
##      concavity_mean      concave.points_mean      symmetry_mean
##      -0.25840048      -0.26085376      -0.13816696
##      fractal_dimension_mean      radius_se      texture_se
##      -0.06436335      -0.20597878      -0.01742803
##      perimeter_se      area_se      smoothness_se
##      -0.21132592      -0.20286964      -0.01453145
##      compactness_se      concavity_se      concave.points_se
##      -0.17039345      -0.15358979      -0.18341740
##      symmetry_se      fractal_dimension_se      radius_worst
##      -0.04249842      -0.10256832      -0.22799663
##      texture_worst      perimeter_worst      area_worst
##      -0.10446933      -0.23663968      -0.22487053
##      smoothness_worst      compactness_worst      concavity_worst
##      -0.12795256      -0.21009588      -0.22876753
##      concave.points_worst      symmetry_worst      fractal_dimension_worst
##      -0.25088597      -0.12290456      -0.13178394
```

#Q10. What is the minimum number of principal components required to explain 80% of the variance of the data? 5 PCs

```
summary(wisc.pr)
```

```
## Importance of components:
##          PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  3.6444 2.3857 1.67867 1.40735 1.28403 1.09880 0.82172
## Proportion of Variance 0.4427 0.1897 0.09393 0.06602 0.05496 0.04025 0.02251
## Cumulative Proportion 0.4427 0.6324 0.72636 0.79239 0.84734 0.88759 0.91010
##          PC8      PC9      PC10     PC11     PC12     PC13     PC14
## Standard deviation  0.69037 0.6457 0.59219 0.5421 0.51104 0.49128 0.39624
## Proportion of Variance 0.01589 0.0139 0.01169 0.0098 0.00871 0.00805 0.00523
## Cumulative Proportion 0.92598 0.9399 0.95157 0.9614 0.97007 0.97812 0.98335
##          PC15     PC16     PC17     PC18     PC19     PC20     PC21
## Standard deviation  0.30681 0.28260 0.24372 0.22939 0.22244 0.17652 0.1731
## Proportion of Variance 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104 0.0010
## Cumulative Proportion 0.98649 0.98915 0.99113 0.99288 0.99453 0.99557 0.9966
##          PC22     PC23     PC24     PC25     PC26     PC27     PC28
## Standard deviation  0.16565 0.15602 0.1344 0.12442 0.09043 0.08307 0.03987
## Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005
## Cumulative Proportion 0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
##          PC29     PC30
## Standard deviation  0.02736 0.01153
## Proportion of Variance 0.00002 0.00000
## Cumulative Proportion 1.00000 1.00000
```


#Hierarchical clustering # Scale the wisc.data data using the “scale()” function

```
data.scaled <- scale(wisc.data)
data.dist <- dist(data.scaled)
wisc.hclust <- hclust(data.dist,method="complete")
```

#Q11. Using the plot() and abline() functions, what is the height at which the clustering model has 4 clusters? around 20

```
plot(wisc.hclust)
abline(h=20, col="red", lty=2)
```

## Cluster Dendrogram



```
data.dist
hclust (*, "complete")
```

#Selecting number of clusters

```
wisc.hclust.clusters <- cutree(wisc.hclust,k=4)
table(wisc.hclust.clusters, diagnosis)
```

```
##           diagnosis
## wisc.hclust.clusters  B  M
##           1  12 165
##           2   2   5
##           3 343  40
##           4   0   2
```

#Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10? 5 is a better cluster match

```
wisc.hclust.clusters <- cutree(wisc.hclust,k=2)
table(wisc.hclust.clusters, diagnosis)
```

```
##           diagnosis
## wisc.hclust.clusters  B  M
##           1 357 210
##           2   0   2
```

```
wisc.hclust.clusters <- cutree(wisc.hclust,k=10)
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B  M
##           1  12 86
##           2   0 59
##           3   0  3
##           4 331 39
##           5   0 20
##           6   2  0
##           7  12  0
##           8   0  2
##           9   0  2
##          10   0  1
```


```
wisc.hclust.clusters <- cutree(wisc.hclust,k=5)
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B  M
##           1  12 165
##           2   0  5
##           3 343 40
##           4   2  0
##           5   0  2
```

#Using different methods #Q13. Which method gives your favorite results for the same data.dist dataset?  
Explain your reasoning. I love ward.D2 best because it gave me fewer clusters but more clear grouping

```
wisc.hclust.single <- hclust(data.dist,method="single")
plot(wisc.hclust.single)
```

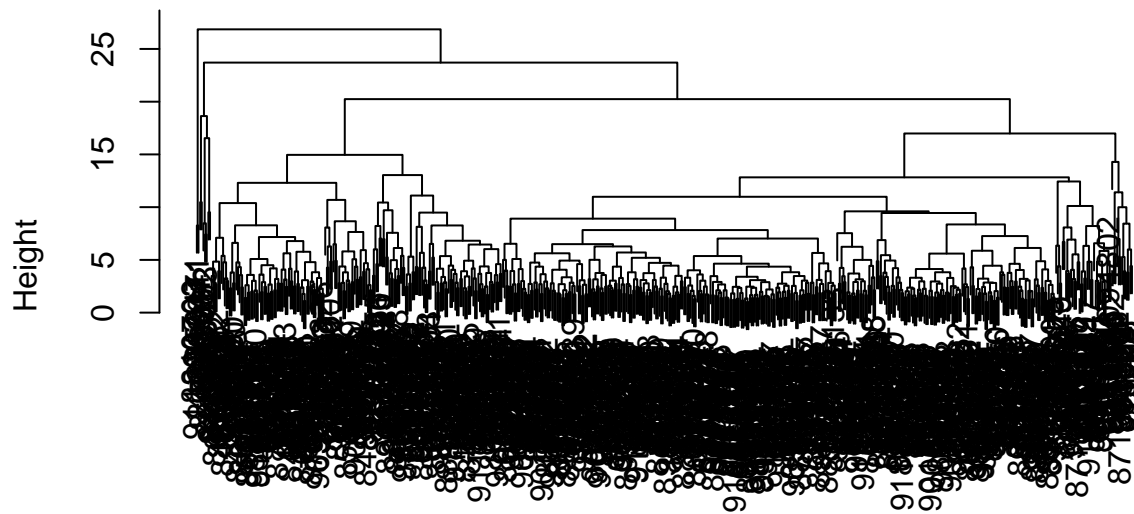
## Cluster Dendrogram



```
data.dist  
hclust (*, "single")
```

```
wisc.hclust.complete <- hclust(data.dist,method="complete")  
plot(wisc.hclust.complete)
```

## Cluster Dendrogram

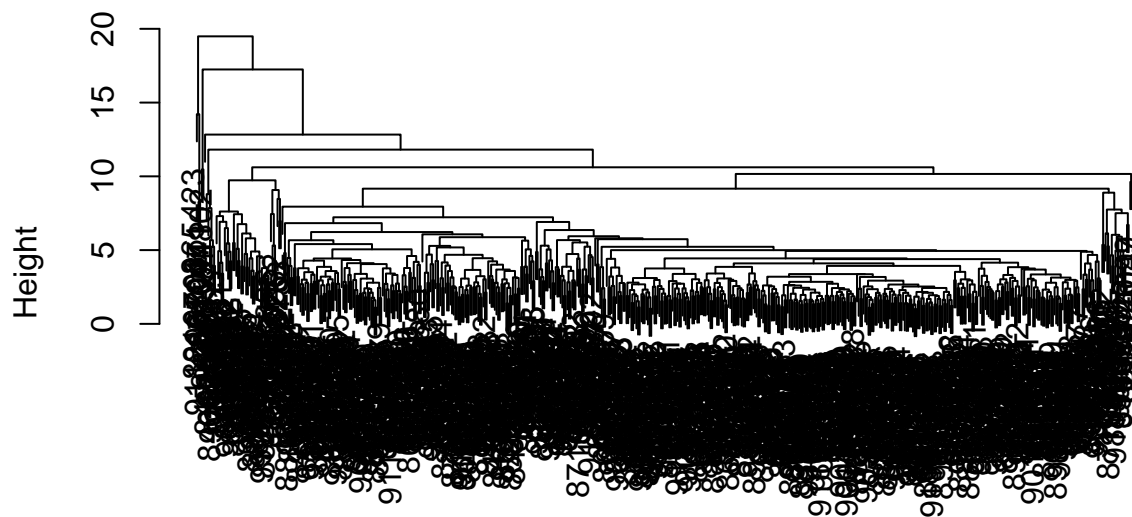


```
data.dist  
hclust (*, "complete")
```

```
wisc.hclust.average <- hclust(data.dist,method="average")  
plot(wisc.hclust.average)
```



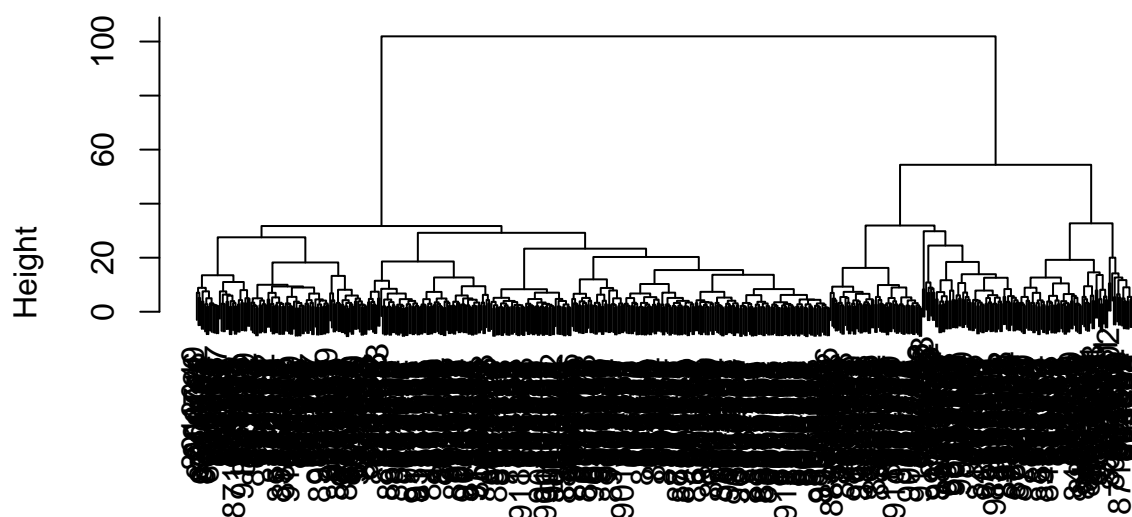
## Cluster Dendrogram



data.dist  
hclust (\*, "average")

```
wisc.hclust.ward <- hclust(data.dist,method="ward.D2")  
plot(wisc.hclust.ward)
```

## Cluster Dendrogram



```
data.dist
hclust (*, "ward.D2")
```

#4. OPTIONAL: K-means clustering #Q14. How well does k-means separate the two diagnoses? How does it compare to your hclust results? Personally, I think k-means separate the two data better since the different clusters are more separated.

```
wisc.km <- kmeans(scale(wisc.data), centers= 2, nstart= 20)
table(wisc.km$cluster, diagnosis)
```

```
##      diagnosis
##      B      M
##  1  14  175
##  2 343   37
```

```
table(wisc.km$cluster, wisc.hclust.clusters)
```

```
##      wisc.hclust.clusters
##      1  2  3  4  5
##  1 160  5  20  2  2
##  2  17  0 363  0  0
```

#Combining methods


```
wisc.pr.hclust <- hclust(dist(wisc.pr$x[,1:7]),method="ward.D2")
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
## grps
## 1 2
## 216 353
```


```
table(grps, diagnosis)
```

```
##      diagnosis
## grps  B   M
## 1  28 188
## 2 329  24
```

```
plot(wisc.pr$x[,1:2], col=grps)
```



```
plot(wisc.pr$x[,1:2], col=diagnosis)
```




```
g <- as.factor(grps)
levels(g)
```

```
## [1] "1" "2"
```

```
g <- relevel(g,2)
levels(g)
```

```
## [1] "2" "1"
```

```
# Plot using our re-ordered factor
plot(wisc.pr$x[,1:2], col=g)
```



#Q15. How well does the newly created model with four clusters separate out the two diagnoses? The newly created model can separate the clusters well

```
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)
table(wisc.pr.hclust.clusters, diagnosis)
```

```
##           diagnosis
## wisc.pr.hclust.clusters  B  M
##           1  28 188
##           2 329  24
```

#Q16. How well do the k-means and hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the table() function to compare the output of each model (wisc.km\$cluster and wisc.hclust.clusters) with the vector containing the actual diagnoses Hierarchical clustering models perform better since it separates clusters into more groups and more separate.

```
table(wisc.km$cluster, diagnosis)
```

```
##      diagnosis
##        B  M
##    1  14 175
##    2 343  37
```

```
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B  M
##              1 12 165
##              2  0   5
##              3 343  40
##              4  2   0
##              5  0   2
```


#Q17. Which of your analysis procedures resulted in a clustering model with the best specificity? How about sensitivity? Variance test aims to increase sensitivity while the clustering process improves specificity.

#Prediction #url <- "new\_samples.csv" #Q18. Which of these new patients should we prioritize for follow up based on your results? Patient 2 needs more follow up

```
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

```
##           PC1      PC2      PC3      PC4      PC5      PC6      PC7
## [1,]  2.576616 -3.135913  1.3990492 -0.7631950  2.781648 -0.8150185 -0.3959098
## [2,] -4.754928 -3.009033 -0.1660946 -0.6052952 -1.140698 -1.2189945  0.8193031
##           PC8      PC9      PC10     PC11     PC12     PC13     PC14
## [1,] -0.2307350 0.1029569 -0.9272861 0.3411457  0.375921 0.1610764 1.187882
## [2,] -0.3307423 0.5281896 -0.4855301 0.7173233 -1.185917 0.5893856 0.303029
##           PC15     PC16     PC17     PC18     PC19     PC20
## [1,] 0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216
## [2,] 0.1299153  0.1448061 -0.40509706  0.06565549  0.25591230 -0.4289500
##           PC21     PC22     PC23     PC24     PC25     PC26
## [1,] 0.1228233 0.09358453 0.08347651 0.1223396 0.02124121 0.078884581
## [2,] -0.1224776 0.01732146 0.06316631 -0.2338618 -0.20755948 -0.009833238
##           PC27     PC28     PC29     PC30
## [1,] 0.220199544 -0.02946023 -0.015620933 0.005269029
## [2,] -0.001134152 0.09638361 0.002795349 -0.019015820
```

```
plot(wisc.pr$x[,1:2], col=g)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



sessionInfo()